



# Operating Systems Lab (C+Unix)

**Enrico Bini**

University of Turin

# Outline

## 1 C: arrays and strings

# Arrays

- An array is **not a class** (as in Java)
- An array is a **contiguous** area of memory allocated to **several variables** of the **same type**
- An array is declared by  
`<type> <identifier>[<size>];`  
it has size `sizeof(<identifier>) = sizeof(<type>)*<size>`
- Example:  
`int v[10];`  
declares the variable `v` as an array of 10 `int` variables.
- Elements are `v[0]`, `v[1]`, ..., `v[<size>-1]` and are stored contiguously in memory
- **C does NOT check array boundaries!** WARNING: `v[-1]` is syntactically correct

| address | content            | variable                  |
|---------|--------------------|---------------------------|
| ...     | ...                | ...                       |
| 0080F8  | ...                | ...                       |
| 0080FC  | <code>v[-1]</code> | outside <code>v!!!</code> |
| 008100  | <code>v[0]</code>  | <code>v</code>            |
| 008104  | <code>v[1]</code>  |                           |
| ...     | ...                |                           |
| 008124  | <code>v[9]</code>  |                           |
| 008128  | <code>v[10]</code> | outside <code>v!!!</code> |
| 00812C  | ...                | ...                       |

# Arrays: length

- The **length** of an array is **not** saved in the data structure
  - ▶ **Do not ever try** to invoke the “method” `length()` with an “array object”
  - ▶ “methods” and “objects” **do not exist** in C
- The programmer must record the length of the array in some way
  - ▶ by storing a special character terminating the useful content (such as in strings, which are terminated by the byte 0)
  - ▶ by recording the length in another (additional) variable
- Still, the following constant expression is useful to compute the length of an array

```
int v[10], len;  /* declaring an array v of a given length */  
  
len = sizeof(v)/sizeof(v[0]);
```

## Strings: arrays of non-zero bytes terminated by 0

- The String “object” or “class” **does not exist** in C
  - ▶ (again, “object” and “class” **do not exist** at all in C)
- In C, the term “string” is used to denote
  - ① an array of char, as declared by:  
`char s[100];`
  - ② the bytes of such an array are interpreted as ASCII codes of characters
  - ③ the byte 0 is written in s after the last character, to terminate the string
- Constant strings are enclosed by **double quotes "**

```
"constant"    /* valid constant string. Allocates 9 bytes */  
'wrong'      /* string are NOT enclosed by single quotes ' */  
'A'          /* ASCII code of 'A' */  
"A"          /* 2-bytes string: 'A', 0 */
```

- A string may be printed by the %s placeholder of the printf as in

```
printf("The string s is \"%s\"\n", s);
```

# Initialization of arrays

- Arrays may be initialized by a sequence of values enclosed within { and }
- ❶ The size of the array may be unspecified and determined by the length of the initialization, as follows

```
int v[] = {1,2,3};
```

- ❷ The following declaration+initialization

```
char v1[] = {'C', 'i', 'a', 'o', 0};
```

```
char v2[] = "Ciao";
```

are equivalent and create an array of **5 bytes** (NOT 4 bytes)  
(strings are arrays of characters terminated by 0)

- ❸ If the size is specified, as in

```
int v[10] = {3, -1, 4};
```

then all following elements are set equal to zero. Hence,

```
int v[100] = {0};
```

is a convenient way to initialize all elements of the vector to zero.

## Strings in memory, converting string into integers

- A string is stored as an array (sequence) of characters, terminated by the null character (0)

```
char v[] = "258";
```

| address          | (hex)  |   |
|------------------|--------|---|
| 7FFF0040671A8108 | 32='2' | v |
| 7FFF0040671A8109 | 35='5' |   |
| 7FFF0040671A810A | 38='8' |   |
| 7FFF0040671A810B | 00     |   |

```
int n = 258; /* =256+2 */
```

| address          | (hex) |   |
|------------------|-------|---|
| 7FFF0040671A8108 | 02    | n |
| 7FFF0040671A8109 | 01    |   |
| 7FFF0040671A810A | 00    |   |
| 7FFF0040671A810B | 00    |   |

- Converting a string into an integer

```
#include <stdlib.h>
```

```
int a;
```

```
a = strtol(s, NULL, 10); /* 10 is the base of conversion */
```

- ▶ stores the value represented by the string s in the integer variable a
- ▶ the second parameter is for advanced users

# Strings in memory, converting string into floating point

- Converting a string into a floating point number

```
#include <stdlib.h>

double a;

a = atof("123.45"); /* same as a = 123.45 */
```

- ▶ stores the value represented by the string *s* in the floating point variable *a*



## Strings: manipulation by including `string.h`

- By including the library `#include <string.h>` some useful function strings to manipulate strings may be used

- ① The following function returns the number of bytes in `s` before the terminating byte 0

```
strlen(s);
```

- ② to append string `src` to string `dest`

```
strcat(dest, src);
```

- ★ `dest` **must** be allocated at least `strlen(dest)+strlen(src)+1`
- ★ otherwise (quoting from `man strcat`): “If `dest` is not large enough, program behavior is unpredictable; **buffer overruns are a favorite avenue for attacking secure programs.**”

- ③ to append **up to** `n` bytes of `src` to string `dest`

```
strncat(dest, src, n);
```

- ★ if no 0 byte terminating `scr` among the first `n` bytes, **only first `n` bytes are concatenated**
- ★ it **prevents the user to write arbitrary-long data**

## Reading input from the keyboard: `fgets()`

- the function `fgets(...)` reads a string of characters
- `#include <stdio.h>` must be added on top to use it
- Syntax

```
char s[80];  
  
fgets(s, sizeof(s), stdin);
```

- ▶ `s[80]` is a pre-allocated array of characters (**string** of characters)
- ▶ reads a string from `stdin` (**standard input**)
- ▶ store the string up to `sizeof(s)-1` characters into `s`. The string **cannot** be `sizeof(s)` long because the terminating zero must be stored too
- ▶ the string is read until EOF (end-of-file, Ctrl+D) or newline
- ▶ if “Enter” is pressed, then the ASCII code of “new line” (=10) is also stored in `s`

- `man fgets`

*test-read.c*, try with input from file