



Operating Systems Lab (C+Unix)

Enrico Bini

University of Turin

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- Variables on the heap
- Variables in memory: comparison
- Variables stored in processor registers

Scope of a variable

- The *scope* of a name (variable or function) is the portion of code where that name is visible and then it may be used
- The scope of a name (variable) declared within a function is restricted to that function
- Parameters of functions have the same scope of a variable declared inside a function
- The scope of a name declared outside any function (*global variables or function declarations*) is from the place of declaration until the end of the file
- If a variable with the same name is declared both outside and inside a function, the one inside the function prevails

Global variables

- Global variables are declared outside any function
- Global variables are visible to all functions from the declaration to the end of file
- Usage of global variables:
 - ▶ good: when many functions share a large amount of data, the usage of global variable is more efficient (it prevents parameter passing)
 - ▶ bad: the code relying much on global variable may be:
 - 1 hardly portable: functions are not really “isolated” pieces of code
 - 2 hard to comprehend/debug: when the reader finds a global variable, it may be not obvious where it is declared
- If a function uses a global variable as input or output, it is **strongly recommended** to add a comment on top of the function
- Names of global variables should be highly informative to avoid the reader to browse much code:
 - ▶ `number_students` is good
 - ▶ `n` is bad


Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- Variables on the heap
- Variables in memory: comparison
- Variables stored in processor registers

Storage classes

- All C variables have a **storage class**, which determines where variables are stored
- Normally, variables are stored in memory. **Three possible areas of memory:**
 - ① variables over the **BSS** (Block Standard by Symbol, historical acronym)
 - ② variables over the **stack segment**
 - ③ variables over the **heap**
- Moreover,
 - ④ **variables may be stored in registers**
- Finally,
 - ⑤ the storage class may be delegated to other places in the code (*external variables*). More details in “Modular programming”

Memory segments

- Depending on the needs, the OS assigns a few memory segments to *processes* (which are programs in execution)
- Segments are mapped over the process address space
- Each segment has:
 - ① start and end addresses (meaningful in the process address space)
 - ② flags that determine the access modes:
 - ★ read: it can be read
 - ★ write: it can be written
 - ★ execute: it contains code which may be executed
 - ★ private/shared: if it isn't/is shared among other processes
- To view the memory mapping of a process, try:
 - `ps -Af | grep sh` to get a Process ID (PID)
 - `cat /proc/<PID>/maps` to print the memory map of <PID>

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- Variables on the heap
- Variables in memory: comparison
- Variables stored in processor registers

Variables on the BSS

- BSS is a read-write memory segment
- Size of BSS is decided at compile time (depending on the size of allocated variables, plus some padding for alignment)
- Two ways to allocate variables over the BSS
 - 1 global variables
 - 2 local variables declared with the static qualifier

```
void func(void) {  
    static int my_static_var;  
    ...  
}
```

- Allocated: at the begin of the program
- Deallocated: at the end of the program

static variables within functions

- Scope: same as local variables (only within the function)
- Lifetime: same as global variables (from the start to the end of the program)
- Typical usage: to keep some state between consecutive invocation of a function

```
void func(void) {  
    static int count_invocations = 0;  
  
    ++count_invocations;  
    ...  
}
```

- Example:
test-static.c

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- **Variables on the stack**
- Variables on the heap
- Variables in memory: comparison
- Variables stored in processor registers

Variables on the stack

- When variables are declared at the top of a function, they are allocated onto the stack (unless the `static` qualifier is pre-fixed)
- Variable are allocated over the stack by reducing the stack pointer as needed by the size of the variables
- Allocated: when the function is entered
- Deallocated: when returning from the function
- Hence, we cannot rely on the their initial value
- The prefix `auto` in variables declaration, such as in

```
void func(void) {  
    auto int my_stack_var;  
    ...  
}
```

may be used. However, since it is the default allocation it is rarely (never?) used explicitly

- How much space is available on the stack?
test-stack-killer.c

Content of the stack

- The stack is a memory area with **LIFO (Last-In First-Out)** policy
 - ▶ push assembly instruction stores data to top of the stack
 - ▶ pop assembly instruction extracts data from the top of the stack
- Main purpose is to store parameters and return address of function invocation
 - ▶ when a function is invoked (`call` assembly instruction),
 - 1 the parameters of the function invocations are pushed to the stack
 - 2 the return address is pushed to the stack
 - 3 then the control flow goes to the invoked function
 - ▶ when we return from function (`ret` assembly instruction))
 - 1 the return address is fetched from the stack
 - 2 the control goes back to the invoking function
- The example *test-stack.c*
 - ▶ shows the content of the stack (the parameters, the return address, etc.)
 - ▶ alters the content of the stack to modify the execution flow (typical attack)

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- **Variables on the heap**
- Variables in memory: comparison
- Variables stored in processor registers

Variables on the heap: dynamic allocation

- The heap (in Italiano “mucchio”, “cumulo”) is a memory area available to the program upon specific request to the operating system
- The program may ask the OS some memory via the following calls

```
#include <stdlib.h>

void * malloc(size_t size);
void * calloc(size_t nmemb, size_t size);
void * realloc(void *ptr, size_t size);
```

which is returned via a pointer (void *)

- ▶ malloc allocates size bytes in memory
- ▶ calloc allocates nmemb elements of size bytes in memory and set them to zero
- ▶ realloc changes the size of previously allocated area
- Allocating memory via malloc is called *dynamic memory allocation* because the size of the allocated memory is **decided at runtime**
- When variables are declared, the size of memory is decided at compile time (*static memory allocation*)

Standard ways for dynamic allocation

- Standard code to allocate an array of `num` elements

```
int * p;  
  
p = malloc(num*sizeof(*p));  
/* better than malloc(num*sizeof(int)) */
```

- `calloc(...)` has a slightly different syntax and it clears the memory (set all bytes equal to zero)

```
int * p;  
  
p = calloc(num, sizeof(*p));
```

- After the allocation, the memory can be used as needed

```
p = calloc(num, sizeof(*p));  
for (i=0; i<num; i++) {  
    p[i] = i*i; /* using array notation */  
}
```


Memory must be freed

- All memory areas allocated by malloc, calloc and realloc must be released by free
- standard code to deallocate a memory area pointed by p is

```
free(p);
```

- free(p) is error, if p not returned by malloc/calloc
- A special care must be taken to free a memory area before the pointer to the area is lost

```
p = malloc(N*sizeof(*p));  
...  
p = &v;  /* ref to allocated mem is lost!! */
```

- To avoid forgetting to free the memory, it is recommended to write the free code immediately, possibly at the bottom of the file.
- Lifetime of memory allocated onto the heap
 - ▶ Allocated: when malloc()/calloc() is invoked
 - ▶ Deallocated: when free() is invoked (or at the end of the program)

Static vs. dynamic allocation

- Is it better static allocation

```
int v[100];
```

- or, dynamic allocation

```
int * v;
```

```
v = malloc(100*sizeof(*v));
```

- used by same syntax: `v[10] = 412;`
- **Dynamic allocation can use less memory** than static allocation (static allocation requires overallocating, by dynamic allocation memory can be allocated when needed)
- **Static allocation is faster** since it avoids expensive system calls such as `malloc` and `free`
- Example of usage of `malloc`
test-malloc.c

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- Variables on the heap
- **Variables in memory: comparison**
- Variables stored in processor registers

Allocation of data in memory

- Let us have a look to the following examples:

- ▶ *test-show-addr.c*

- Remember the difference between

```
char v[] = "string0";  
char * p = "string1";
```

- ▶ string0 may be modified (it belongs to a page with “w” permission)
 - ▶ string1 may not be modified (no “w” permission)

Outline

1 C: scope of variables

2 C: storage classes

- Variables on the BSS
- Variables on the stack
- Variables on the heap
- Variables in memory: comparison
- Variables stored in processor registers

register variables

- The compiler may be informed that some variable **should be allocated to a register** of the processor by adding the keyword `register` at the declaration

```
register int my_register_var;
```

- register variables are used for frequently accessed variables: access time to a register is **10–100 times faster** than access to memory
- The number of register **is limited**: the compiler cannot guarantee the allocation to a register