



# Operating Systems Lab (C+Unix)

**Enrico Bini**

University of Turin

# Outline

## 1 Files

- Streams
- File descriptors

# Files and file descriptors

Two different interfaces to files

## ① *streams* of type

```
FILE * my_f;
```

(FILE is a struct defined in `stdio.h`) and

- ▶ input/output is *buffered* to improve performance:
  - ★ inconvenient to write on a disk byte by byte
  - ★ data to be written to disk is stored to a memory area (*buffer*) only
  - ★ the buffer is written to the disk (*flushed*) depending on the buffering policy

## ② *file descriptors* of type

```
int my_fd;
```

a file descriptor is just an integer (which is the index in a table managed by the operating system)

- ▶ lower level interface
- ▶ not buffered
- ▶ file descriptors are used for more general purpose than writing on a disk file (interprocess communication, communicate via TCP/IP, etc.)

# Outline

## 1 Files

- Streams
- File descriptors

# Streams: opening/closing

- Before being used streams must be opened by

```
FILE * fopen(const char *path, const char *mode);
```

- ▶ path is the path of the file to be open
- ▶ mode is a **string** (not a character) specifying the read/write opening mode. Example: "rw" Check: **man fopen** for full description
- ▶ a pointer (FILE \*) is returned

Example: opening "my\_file.txt" in read mode

```
FILE * my_f;  
  
my_f = fopen("my_file.txt", "r");
```

- After its usage, a stream must be closed by

```
int fclose(FILE * stream);
```

if streams are not closed, the OS may be unable to open new files

## Streams: reading

- For each open file, the OS keeps and updates a *file position indicator*
- Every reading happens at the current position, which is incremented by the number of bytes read

```
int fgetc(FILE *stream);  
char *fgets(char *s, int size, FILE *stream);  
int fscanf(FILE *stream, const char *format, ...);
```

- they all read from the current position
- fgetc reads and returns the byte (char) read into a (int). If end-of-file is reached, then the int non char-representable EOF macro is returned (typically of value -1)
  - ▶ the byte 0xFF can be distinguished by EOF
- fgets reads an array of up to size bytes. Returned NULL if end-of-file is reached.
- fscanf read by scanf/printf format

## Streams: writing

- Every writing happens at the current position, which is incremented by the number of written bytes

```
int fputc(int c, FILE *stream);  
int fputs(const char *s, FILE *stream);  
int fprintf(FILE *stream, const char *format, ...);
```

- fputc writes the byte c, casted to char, to the file
- fputs writes the zero-terminated string s **without** the terminating 0 byte
- fprintf writes to file by the printf format

# Streams: controlling the position over a file

- The position over a file can be controlled by `fseek()`

```
int fseek(FILE *stream, long offset, int whence);
```

sets the file pointer of `stream` as follows:

- ▶ if (`whence == SEEK_SET`), position is set equal to `offset`
- ▶ if (`whence == SEEK_CUR`), position is moved by `offset`
- ▶ if (`whence == SEEK_END`), position is moved by `offset` from the end

notice that `offset` may be negative (to move the position backward)

- To know the current position over a file

```
long ftell(FILE *stream);
```

- ▶ The first byte of a file is at position `0`
- ▶ The last byte of a file is at position `<size>-1`
- ▶ When the position is equal to `<size>`, then we reached the end-of-file

- *test-file.c*



## Standard streams: `stdin`, `stdout`, `stderr`

- `stdin`, `stdout`, and `stderr` are all streams (of type `FILE *`) defined by the operating system with a special usage
- `stdin` is “standard input” and it is the stream of characters entered by the keyboard
- `stdout` is “standard output” and it is the stream of characters printed on the terminal
- `stderr` is the “standard error” stream. It is used to print error messages and it is printed on the terminal as well

# Streams: buffering

- The interaction between the (fast) processor and the (slow) devices may degrade the performance
  - ▶ it is not convenient to write a single byte to the disk every time `fputc()` is invoked
- I/O may be buffered: “buffered” read/write are delayed until the “buffering” condition is true. Three types of buffering
  - ① **unbuffered**: all I/O operations happen immediately
  - ② **block buffered**: I/O operations are executed when the buffer is full
  - ③ **line buffered**: I/O operations are executed when newline ‘`\n`’ read
- **stdout is line-buffered**
- **stderr is not-buffered** (normally, we want to see the error messages as soon as they happen): during debugging use `stderr`
- **other files are block buffered, unless specified differently**

# Streams: controlling the buffering

- To force the buffer to be written to the device

```
int fflush(FILE *stream);
```

by `fflush(NULL)`, all open output streams are flushed.

- the function `setvbuf` changes the buffering policy of stream

```
int setvbuf(FILE *stream, char *buf, int mode,  
            size_t size);
```

if mode is:

- ① `_IONBF`, stream is unbuffered (every single byte is written/read immediately)
- ② `_IOLBF`, the buffer is written as soon as newline is found
- ③ `_IOFBF`, write to disk only when buffer is full

**man setvbuf** for more information

- Examples

```
/* set no buffering to stream */  
setvbuf(stream, NULL, _IONBF, 0);
```

# Outline

## 1 Files

- Streams
- File descriptors

# File descriptors and files

- **Streams** (not files) are of type `(FILE *)` (the name “FILE” is only for historical reasons)
- **File descriptors** are of type `int` (sometime called “I/O streams”)
- A file descriptor (fd) identifies a **source/destination of a sequence of bytes**
- File descriptors are a lower level interface than streams
- File descriptors are more general than streams
  - ▶ **all streams have a file descriptor**
  - ▶ there may be file descriptors which are not streams
- File descriptors are opened by different functions depending on their usage:
  - ▶ `int open(...)` (not `fopen(...)`) binds a file in the file system to the returned descriptor
  - ▶ `int socket(...)` binds the data coming-from/going-to a UDP/TCP (and others) connection to the returned descriptor
  - ▶ `pipe(...)` creates a “pipe”: two descriptors attached to each other (more details later in the course)

# File descriptors linked to standard streams

- `stdin`, `stdout`, and `stderr` are standard streams opened by the OS and allowing the program to:
  - ▶ read from keyboard (from `stdin`)
  - ▶ write normal output to terminal (to `stdout`)
  - ▶ write error messages to terminal (to `stderr`)
- Standard file descriptors are associated to these streams:
  - ▶ the integer 0 is the file descriptor of `stdin`
  - ▶ the integer 1 is the file descriptor of `stdout`
  - ▶ the integer 2 is the file descriptor of `stderr`

## Redirecting stdout and/or stderr

- To redirect stdout to a file, truncate if existing  
`COMMAND 1> filename`
- To redirect stdout to a file, append if existing  
`COMMAND 1>> filename`
- To redirect stderr to a file, truncate if existing  
`COMMAND 2> filename`
- To redirect stderr to a file, append if existing  
`COMMAND 2>> filename`
- To redirect stdout and stderr to a file, truncate if existing  
`COMMAND &> filename`
- To redirect stdout and stderr to a file, append if existing  
`COMMAND &>> filename`

# Opening/closing a file descriptor of a file

```
int open(const char *pathname, int flags);
```

- `open(...)` opens a file and returns a fd (**man 2 open** for details)
  - ▶ `pathname`, a string with the pathname of the file
  - ▶ `flags`, specifies how to open (about 20 flags).  
Flags are set by making the bitwise OR “|” among the selected macros
  - ▶ Each macro has one “1” bit only
    - ★ **must include** one among `O_RDONLY`, `O_WRONLY`, `O_RDWR`
    - ★ `O_APPEND`, file opened in append mode
    - ★ `O_CREAT`, create the file if doesn't exist
    - ★ `O_TRUNC`, if file exists, it is truncated
- After being used, file descriptors must be closed

```
int close(int fd);
```

otherwise we may run out of available file descriptors



## Reading from a file descriptor

```
ssize_t read(int fd, void *buf, size_t size);
```

- reads from the file descriptor `fd` up to `size` bytes and store them to `buf`
- it returns the number of bytes actually read (it may be less than `size`)
- if it returns zero, then end-of-file is reached
- if it returns `-1` then an error has occurred

# Writing to a file descriptor

```
ssize_t write(int fd, const void *buf, size_t size);
```

- write `size` bytes from the buffer `buf` to the file descriptor `fd`
- it writes immediately the data, not buffered as `fprintf`
- formatted output over a file descriptor `fd` by

```
int dprintf(int fd, const char *format, ...);
```

- WARNING: by mixing `fprintf` and `write` to the same `fd`/stream you must be careful
  - ▶ `fprintf` uses a buffer to write, while `write` doesn't
  - ▶ the output written by `fprintf` may be delayed w.r.t. the output made via `write`
- *test-buf.c*

# Positioning over a file descriptor

- This position over a file descriptor is controlled by `lseek()`

```
off_t lseek(int fd, off_t offset, int whence);
```

- set the file pointer of `fd` as follows:
  - ▶ if (`whence == SEEK_SET`), position is set equal to `offset`
  - ▶ if (`whence == SEEK_CUR`), position is moved by `offset`
  - ▶ if (`whence == SEEK_END`), position is moved by `offset` from the end
- notice that `offset` may be negative (to move the position backward)
- File descriptors of different types (not associated to files) do not allow positioning by `lseek(...)`