



# Operating Systems Lab (C+Unix)

**Enrico Bini**

University of Turin

-

# Outline

## 1 Signals

- Sending signals
- Handling signals
- Lifecycle of signals: delivering, masking, merging
- Getting a signal when in waiting state

# Signals

- **Signals** are software interrupts delivered to processes.
- Signals can be generated by **user, software, or hardware events**
- Example of signals are:
  - ▶ SIGFPE “Floating Point Exceptions” such as division by zero
  - ▶ SIGILL trying to execute an “Illegal instruction”
  - ▶ SIGINT used to cause program interrupt (Ctrl+C)
  - ▶ SIGKILL causes immediate program termination
  - ▶ SIGTERM polite version of terminating a program (SIGTERM can be handled by the user)
  - ▶ SIGALRM received when a timer (set by `alarm(int seconds)`, `timer_create(...)`, or other calls) has expired
  - ▶ SIGCHLD sent to a parent when a child terminates
  - ▶ SIGSTOP/SIGCONT stop/continue a process
  - ▶ SIGUSR1/SIGUSR2 user-defined signals

`man 7 signal` for a full list

# Outline

## 1 Signals

- Sending signals
- Handling signals
- Lifecycle of signals: delivering, masking, merging
- Getting a signal when in waiting state

# Sending a signal to any process

## ① From a C program

```
#include <sys/types.h>
#include <signal.h>

int kill(pid_t pid, int signum);
```

- ▶ signum, the ID of the signal
  - ★ by sending the signal 0 (*null signal*) we can test the existence of a pid
    - ① if (errno==ESRCH), pid doesn't exist
    - ② if (errno==EPERM), pid exists but no permission to send signals
    - ③ if successful, pid exists and we can send signals
- ▶ pid, the target process
  - ★ if -1 the signal is sent to all (allowed) processes

## ② From command line: kill, signal, and PID

```
kill -INT <PID>
```

```
kill -SIGINT <PID>
```

```
kill -2 <PID>
```

- ▶ If no signal is specified, then SIGTERM is the default
- ▶ `sudo kill -9 -1` is an interesting experience...

# Sending a signal to myself

## ① `raise(signum)`

```
#include <signal.h>

int raise(int signum);
```

to send `signal signum` to myself now.

## ② `alarm(int sec)`

```
#include <unistd.h>

unsigned int alarm(unsigned int sec);
```

asks the OS to send me `SIGALRM` after `sec` seconds

- ▶ returns the seconds remaining until any previously scheduled alarm was due to be delivered, or zero if there was no previously scheduled alarm

## ③ we can also set a periodic timer

```
#include <sys/time.h>

int setitimer(int which, const struct itimerval *new_value,
              struct itimerval *old_value);
```

**man setitimer**

# Outline

## 1 Signals

- Sending signals
- **Handling signals**
- Lifecycle of signals: delivering, masking, merging
- Getting a signal when in waiting state

# Signal handler: default action

Each signal has a **default handler** (`man 7 signal`)

- **Term**, terminate the process.
- **Ign**, ignore the signal.
- **Core**, terminate the process and dump “core”. The core dump file contains the image of the process memory at the time of termination and can be used by a debugger (such as `gdb`) to inspect the causes of termination
- **Stop**, stop the process
- **Cont**, continue the process if it is currently stopped.



# Signal handler: user-defined action

- The **user-defined signal handler is a function that is called asynchronously** at the time of signal delivery, wherever in the code this happens
- At the time of signal delivery
  - 1 the state of the process is **saved** (registers, etc.)
  - 2 the **function of the handler** is executed and then returned
  - 3 the state of the process is **restored**
- The signal handler is very similar to an interrupt handler
- Some **signals (example: SIGKILL) cannot be handled by the user.** Otherwise, an immortal process may be allowed by handling SIGKILL
- **synchronous** signal delivery is possible
  - ▶ the process may want to know the precise moment of signal delivery and wait for it, if needed
    - ★ **sigwaitinfo(...), sigtimedwait(...), and sigwait(...)** suspends the execution until a signal is received

out of the scope of this course

# Signal handler: definition of sigaction

- The user may set a signal handler by the sigaction() system call

```
#define _GNU_SOURCE /* necessary from now on */
#include <signal.h>

int sigaction(int signum,
               const struct sigaction *act,
               struct sigaction *oldact);
```

- 1 **signum**, the number of the signal to be handled
  - 2 **act**, new handler of the signal, if NULL handler unchanged
  - 3 **oldact**, pointer to the old handler, if NULL no handler returned
- **WARNING: sigaction is both a sys call and a struct**

```
sigaction(signum,&new,NULL); /*set new handler*/
sigaction(signum,NULL,&old); /*get cur handler*/
sigaction(signum,&new,&old); /* do both */
```

# Format of the sigaction structure

- Signal handlers are specified by a sigaction data structure

man 2 sigaction

```
struct sigaction {  
    void      (*sa_handler)(int signum);  
    sigset_t  sa_mask; /* illustrated later */  
    int       sa_flags; /* illustrated later */  
    /* plus others (for advanced users) */  
};
```

- sa\_handler is a pointer to a function declared as

```
void signal_handler(int signum);
```

- If standard behavior is required, all bytes of “other fields” must be set to 0

*test-signal-handle.c*

## User-defined signal handlers 1/2

- 1 user-defined signal handlers must be attached to the corresponding signal **before** the signal may be released (for example, if SIGALRM is going to be handled, first attach the handler to the signal, then invoke `alarm(...)`)
- 2 often a single function handles many signal and a `switch(signal)/case` selects the proper action for the signal

```
void handle_signal(int signal) {  
    /* signal signal triggered the handler */  
    switch (signal) {  
        case SIGINT:  
            /* handle SIGINT */  
            break;  
        case SIGALRM:  
            /* handle SIGALRM */  
            break;  
        /* other signals */ } }  
}
```

## User-defined signal handlers 2/2

- ① user-defined signal handlers of parents are inherited by child processes
- ② global variables are visible:
  - ① both in the handler code (executed asynchronously upon the reception of a signal)
  - ② and in the rest of the code (executed according to the “normal” flow of the program)
    - ▶ good: global variables offer a way to inform the “main” program of the occurrence of signals
    - ▶ bad: special care must be taken when invoking functions that use global variables

# Global variables and signal handlers

- some functions (including `printf(...)`) use global data structure. The asynchronous arrival of signals may corrupt this data and produce unexpected behavior
  - ▶ `man signal-safety`  
“If a signal interrupts the execution of an unsafe function, and handler calls an unsafe function, then the behavior of the program is undefined”
- functions that can be safely used within a handler are marked by “AS-Safe” (Asynchronous Signal-Safe) in the GNU libc documentation:
  - ▶ `write(...)` is AS-Safe,
  - ▶ `printf(...)` is **AS-Unsafe**, because it uses global variables (the output buffer)
- *test-printf-handler.c*
- Every time you use any library function in a signal handler **check** that it is **AS-safe** at GNU libc documentation
- `errno` is a global variable, which may be overwritten within the signal (if any function writing to `errno` is used). It is recommended to save it and restore it at the end of the handler

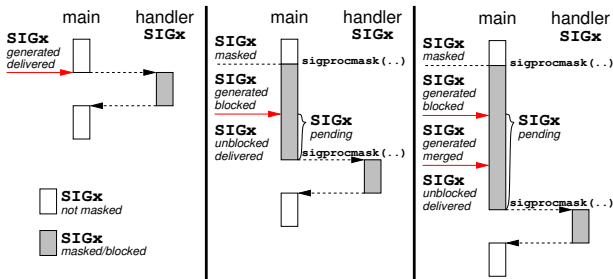
# Outline

## 1 Signals

- Sending signals
- Handling signals
- Lifecycle of signals: delivering, masking, merging
- Getting a signal when in waiting state

# Signals lifecycle

- Signals are *generated* by hardware or software events and are sent to a process
- ① Any process may set a *signal mask* to postpone signal handling
- ② if a signal arrives to a process when it is masked, the the signal is *blocked*
- ③ if the signal is blocked, then it remains *pending*, until it is *unblocked*. As soon as unblocked, it is immediately *delivered* to the process,
- ④ if a signal is generated again while it is already pending, then it is *merged*: after unblocked, it will be delivered only **once**!!
- ⑤ if the signal is not blocked, it is immediately *delivered* to the the process.





# Setting signal masks

- During its execution, each process has its own *signal mask*
- a child process inherits the parent's signal mask
- The signal mask is the collection of signals that are currently *blocked*
- The signal mask of a process can be updated by the system call `sigprocmask(...)` (details later)
- Signal masks are of type `sigset_t`. Functions to manipulate sets:

```
int sigemptyset(sigset_t *set);
int sigfillset(sigset_t *set);
int sigaddset(sigset_t *set, int signum);
int sigdelset(sigset_t *set, int signum);
int sigismember(const sigset_t *set, int signum);
```

`man sigsetops` for more details

# Signal mask of a process

- `sigprocmask(...)` is used to set the signal mask of a process

```
#include <signal.h>

int sigprocmask(int how, const sigset_t *set,
                sigset_t *oldset);
```

- `oldset` is the old mask
- the new mask is set according to:
  - ▶ if (`how==SIG_BLOCK`), then signals in `set` are added to the current mask
  - ▶ if (`how==SIG_UNBLOCK`), then signals in `set` are removed from the current mask
  - ▶ if (`how==SIG_SETMASK`), then `set` becomes the new signal mask
- *test-signal-mask.c*

## Signal masks: why

- Signals arrive asynchronously and may interrupt the program execution at any time
- The programmer must take special care in preventing signals to interrupt the execution in places that leave the status in an inconsistent status  
*test-signal-non-atomic.c*
- type `sig_atomic_t` is an integer and it is guaranteed by the compiler to be accessed atomically (by a single assembly instruction)
- In practice, we can assume that
  - ▶ `int` and
  - ▶ `pointers`are atomic
- To mask or not to mask signals?
  - 1 to mask to avoid inconsistent data
  - 2 not to mask to increase responsiveness and reduce the risk of signal merge

# Signal mask during a handler


- The signal that triggered the handler is masked during the handler
  - ▶ unless the flag `SA_NODEFER` is set (to be explained later)
- `sa_mask` field of `sigaction` sets the mask during the handler
  - ▶ when the handler returns, the set of blocked signals is restored to the value before its execution, regardless of any manipulation of the blocked signals made in the handler

```
/* How to mask a signal during SIGINT handler */
struct sigaction sa;
sigset_t  my_mask;

bzero(&sa, sizeof(sa));          /* clean sa struct */
sa.sa_handler = handle_signal;    /* set handler */
sigemptyset(&my_mask);           /* Set an empty mask */
/* Add a signal to the sa_mask field struct sa */
sigaddset(&my_mask, signal_to_mask_in_handler);
sa.sa_mask = my_mask;
/* Set the handler */
sigaction(SIGINT, &sa, NULL);
```

## Merged signals in handler

- If a signal is generated while it is still pending for being handled, the newly generated and the pending one are merged into one
  - ▶ the presence of a pending signal is stored by a flag only, not by a number (of pending signals)
  - ▶ a signal handler cannot be reliably used to count the number of collected signals!
- *test-signal-merge.c*

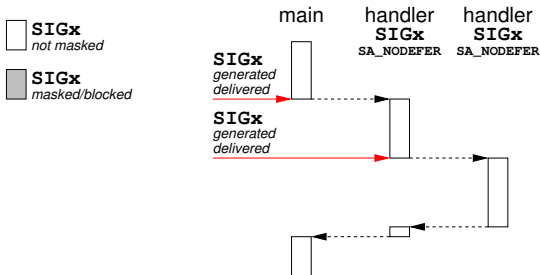


# Unblocking the delivered signal during its own handler

- When a signal `signum` is delivered to a process, during its handler, the signal `signum` is automatically blocked until the handler returns
- The default behavior may be changed by setting the `SA_NODEFER` flag of the struct `sigaction`

```
bzero(&sa, sizeof(sa));  
sa.sa_handler = handle_signal;  
sa.sa_flags = SA_NODEFER; /* nested signals */  
sigaction(SIGUSR1, &sa, NULL);
```

- *test-signal-nodefer.c*



# Outline

## 1 Signals

- Sending signals
- Handling signals
- Lifecycle of signals: delivering, masking, merging
- Getting a signal when in waiting state

# Entering the waiting state: pausing, sleeping

## ① pause()

```
#include <unistd.h>

int pause(void);
```

the process stays in *waiting* state until a signal is caught

## ② sleep(sec) (deprecated)

```
#include <unistd.h>

unsigned int sleep(unsigned int seconds);
```

the process sits in *waiting* state for *sec* seconds. If the process caught a signal while sleeping, then `sleep` returns the remaining second to sleep

Deprecated because from **man 3 sleep**: “`sleep()` may be implemented using `SIGALRM`; mixing calls to `alarm()` and `sleep()` is a **bad idea**.”



# Sleeping by nanosleep

- if sleeping for some time is needed, nanosleep is the right alternative to sleep

```
#include <time.h>

struct timespec my_time;

my_time.tv_sec = 1;
my_time.tv_nsec = 234567000;
nanosleep(&my_time, NULL);
```

sleeps for 1.234567 seconds

**man nanosleep**

## “Synchronization” by sleep/nanosleep: **WRONG**


- Scenario: parent process must wait that child completes some work
- Here's a tempting (wrong) code to synchronize the two processes

```
...  
if (fork()) {  
    /* PARENT */  
    sleep(10);  
    /* the parent thinks the child has finished */  
} else {  
    /* CHILD */  
    /* do my work before the parent checks */  
}
```

- why wrong?
  - 1 we don't know if the child will take less than 10 seconds
  - 2 the OS may decide not to schedule the child for more than 10 seconds
- sleep(sec) only guarantees that the process sleeps for sec
- sleep(...)/nanosleep(...) are used mostly to show output slowly to the user
- **Never** use sleep(...)/nanosleep(...) to wait for another process

# Delivery of signals to a *waiting* process

- Signal handler executes asynchronously:
  - ① the state of the process is saved (registers, etc.)
  - ② the function of the handler is executed
  - ③ the state of the process is restored
- What happens when a signal is delivered to a *waiting* process?
  - ▶ “waiting process”: process waiting on `wait()`, on `pause()`, or `sleep()`, etc... Counted as “sleeping” in `top`
  - ① the process is not executing, since it is waiting on some system call
    - ★ its state is already saved
  - ② the function of the handler is normally executed
  - ③ upon the return of the handler, **two** possible behaviors:
    - ★ **restarting**: the system call is restarted when the handler returns, **OR**
    - ★ **aborting**: the system call aborts, `errno` is set to `EINTR`



## Signals when waiting: selecting the desired behavior

- Which behaviour among “**restarting**” or “**aborting**”?
- The default is “**aborting**”
  - ① the waiting system call returns an erroneous value: -1, etc. depending on the call
  - ② `errno` is set equal to `EINTR`
- If the “**restarting**” behavior is desired, then consider
  - ① setting the flag `SA_RESTART` in the `sa_flags` field of the struct `sigaction`
  - ② checking if `(errno == EINTR)` after the waiting call and possibly re-invoke the call
- Unfortunately, different calls have different behavior
- It is then recommended to check the full documentation at `man 7 signal`

Section “Interruption of system calls and library functions by signal handlers”
- *test-signal-when-wait.c*