

# Operating Systems Lab (C+Unix)

**Enrico Bini**

University of Turin

# Outline

1 SysV: IPC

2 SysV: message queues

# Inter-Process Communication (IPC)

- Processes may communicate via *IPC objects*
  - ▶ *message queues* allow processes to send and receive messages
  - ▶ *shared memory* allows processes to view a common area of memory where all processes can write/read
  - ▶ *semaphores* enable only a few process to access a shared resource or enable synchronization
- IPC objects are implemented by (at least) two standards:
  - 1 *System V*, older standard: first released by AT&T in 1983,
  - 2 *POSIX*, more recent standard inspired by System V, rapidly spreading and adopted by many

They are both available in many Unix-systems, such as Linux

- The course will adopt System V standard, in the future we may switch to POSIX
- Some documentation can be found at <http://www.tldp.org/LDP/lpg/node7.html>  
(Warning: it is dated 1995!)

# IPC objects: persistence

- IPC objects are *persistent*: they survive in the kernel space even after all the processes (creating or accessing the object) have terminated
  - ▶ this is good: IPC objects enable the communication between
    - 1 processes that just know the “name” of the IPC object (such as in FIFOs)
    - 2 even different invocations of the same executable
  - ▶ this is bad: it is worse than forgetting to free a dynamically allocated memory (by malloc)
    - ★ if not explicitly removed (when needed), they can quickly fill up the memory
- It is possible to create, list or erase current IPC objects at command line
  - ▶ the command `ipcs` shows the status of current IPC objects
  - ▶ the command `ipcs -l` shows the system limits on the resources
    - ★ also available in the `/proc/sys/kernel/` file system
    - ★ can be modified by `sudo sysctl kernel.msgmax=65536`
  - ▶ the command `ipcmk` creates an IPC object (only Linux, not standard)
  - ▶ the command `ipcrm` erases the specified IPC object

# IPC objects: IDs and keys

- Any System V IPC *object* (message queue, shared memory, or semaphore) is identified by a unique identifier (ID) of type `int`
  - uniqueness is per type: there may be two IPC objects of different type with the same ID
- Processes that are willing to use the same IPC object for communication, must both know its ID
  - Processes may get the ID from a common *key*
- For each type of IPC object the `get()` function returns the ID from a key
  - `int msgget(key_t key, ...)` to get a message queue
  - `int semget(key_t key, ...)` to get a semaphore
  - `int shmget(key_t key, ...)` to get a shared memory
- How can two processes agree on a key?
  - the key can be hard-coded in a common `.h` file (via `#define`)
  - the key can be `IPC_PRIVATE` to create a new object (not really private since it may be shared, unfortunate choice of name)
  - the key can be `getppid()`, if object shared among siblings

## Getting an IPC object from a key

- All three types of IPC objects have a similar method to get the ID
- Any process accessing the object should call the `???get()` functions to get the ID, **unless the ID is already known**:
  - ▶ **inherited from parent by `fork()`**
  - ▶ **passed as parameters at invocation time**

```
int msgget(key_t key, int flags);  
int shmget(key_t key, size_t size, int flags);  
int semget(key_t key, int nsems, int flags);
```

- the IPC object identifier associated to key is returned. It may be:
  - ▶ the ID of a new object just created by calling `???get()`
  - ▶ the ID of an existing object previously created by others

Check next slide for the precise behaviour

- **flags specifies the read/write permissions** of user/group/others in the standard octal form
  - ▶ 0400 read to user
  - ▶ 0020 write to group
  - ▶ 0666 read/write to everybody
  - ▶ ...
- Also flags may include macros in bitwise OR

## Four ways to “get” an IPC object (Example: msg queues)

- 1 Create a **new object** and **communicate the ID to processes**

```
id = msgget(IPC_PRIVATE, 0600); /* 0600 is an example */
```

- ▶ then give the ID via command-line (through `execve(...)`) or
- ▶ use the copied variable through `fork()`

- 2 Create a **new object** with a **given** key

```
id = msgget(key, IPC_CREAT | IPC_EXCL | flags);
```

- ▶ if IPC object with key **exists**, return -1 and `errno=EEXIST`
- ▶ if IPC object with key does **not exist**, it is created

- 3 Use an **existing object** with a **given** key

```
id = msgget(key, 0644); /* 0644 is an example */
```

- ▶ the ID of the **existing** object associated to key is returned
- ▶ -1 returned and `errno=ENOENT` if no IPC object exists with key

- 4 Use an **existing object** with a **given** key, **or create it if not exists**

```
id = msgget(key, IPC_CREAT | 0660);
```

- ▶ if IPC object with key **exists**, same as `msgget(key, flags)`
- ▶ if IPC object with key does **not exist**, same as  
`id = msgget(key, IPC_CREAT | IPC_EXCL | flags);`

# Typical issues due to persistence

- 1 Say that your program creates and uses some IPC object
- 2 Say that your program crashes or it never ends and you have to stop it by Ctrl+C
- 3 Then you fix it and you launch it again

**The IPC object of the second run still has the same content it had after the first run!!!**

- Possible fixes:
  - 1 “get” the object with `IPC_PRIVATE` key it always returns a new object
    - ★ may run out of memory
  - 2 install Ctrl+C handler that cleans up objects
  - 3 remove old objects at command line by `ipcrm`



# Outline

1 SysV: IPC

2 SysV: message queues

# Lifecycle of a message queue

- 1 A message queue  $Q$  is created by process  $A$
- 2  $Q$  is opened for being used (send/receive) by processes  $P_1, \dots, P_n$
- 3 Processes  $P_1, \dots, P_n$  send and receive messages over  $Q$  as needed
  - ▶ sent messages are enqueued to the tail
  - ▶ received messages are searched from the head (may pick messages other than the first one)
- 4 Sender processes send messages even if nobody will ever receive them
  - ▶ no SIGPIPE-like method when all read ends are closed (as in pipes)
  - ▶ if queue is full, **a process may be blocked forever**
- 5 Receiver processes cannot know when senders have finished
  - ▶ no EOF-like method when all write ends are closed (as in pipes)
  - ▶ if nobody is sending, **a process may be blocked forever**
- 6 Once sender processes have finished sending their messages, the message queue will persist in the kernel and receiver processes remain blocked waiting for a message until the queue exists
- 7 It is necessary to determine correctly the condition that allows the deallocation of the message queue

# Creating a message queue (or getting ID of an existing one)

- The system call

```
int msgget(key_t key, int msgflag);
```

returns the identifier of a message queue associated to key

- ▶ msgflag is a list of ORed ("|") options including:
  - ★ read/write permissions (least significant 9 bits) in standard octal form (the "execute" (x) permission is ignored)
  - ★ IPC\_CREAT: if queue exists, its ID is returned; if it doesn't exist, it is created
  - ★ IPC\_EXCL (used only with IPC\_CREAT): the call fails (with `errno=EEXIST`) if the queue exists

- Message queues are persistent objects: they will survive to the death of the creator, they must be erased explicitly

# Message format

- Messages **must** start with a long value: the type of a message
  - ▶ the type **must** be strictly positive (not zero): sending a message with non-positive type results in an error
  - ▶ the type can be used to select messages to be read
- For example, the default message structure

```
struct msgbuf {  
    long mtype;           /* type of message */  
    /* my personal message goes here */  
};
```

- The user can define any message structure as long as:
  - 1 the first sizeof(long) bytes are reserved to the message type
  - 2 the total length of the message does not exceed the maximum  
cat /proc/sys/kernel/msgmax
- Messages of length 0 are also acceptable. If so only sizeof(long) bytes are sent
- Do not use pointers in a message: pointers live into the memory of a process. A pointer written by another process does not make sense

# Sending a message to a queue

- Messages are sent by the `msgsnd()` system call

```
int msgsnd(int msqid, const void *msgp, size_t msgsz, int
    msgflg);
```

- The caller process must **have write permissions** on the queue to send a message
  - ▶ **msqid**, the ID of the message queue where the message is sent
  - ▶ **msgp**, pointer to the message structure
  - ▶ **msgsz**, size of the message content (excluding `sizeof(long)` bytes of the heading type)
- If queue is full
  - ▶ the call **`msgsnd()` blocks until some space for the message is made**, or
  - ▶ if flag **`IPC_NOWAIT` is set, it returns `-1` with `errno = EAGAIN`**
- After processes have finished sending, they cannot close their “write end” as in pipes
  - ▶ Message queues are “closed” (erased) separately once they are no longer needed

## Receiving a message

- To receive a message from the queue `msqid` and copy it to the buffer pointed by `msgp` the `msgrcv()` system call is used

```
int msgrcv(int msqid, void *msgp,  
           int msgsz, long mtype, int msgflg);
```

- Process must have **read permissions on the queue to receive a msg**
- `msgsz` is the size of the message (without type) copied to the buffer
- The received message is selected as follows:
  - ▶ if (`mtype == 0`), the first message in the queue is selected
  - ▶ if (`mtype > 0`), the first message of type `mtype` is selected
  - ▶ if (`mtype > 0`) and `MSG_EXCEPT` flag is set, the first message of type **different than** `mtype` is selected
  - ▶ if (`mtype < 0`), the first message in the queue of the **lowest** type less than or equal to `mtype` is selected (low types have a high priority)
- If no message is selected by the rules above
  - ▶ the call `msgrcv()` blocks until a selected message arrives, or
  - ▶ if flag `IPC_NOWAIT` is set, it returns `-1` with `errno = ENOMSG`
- the received message is erased from the queue (unless `MSG_COPY` flag)

# Errors on sending/receiving messages

- Both `msgsnd()` and `msgrcv()` may fail and return `-1`
- The error code `errno` is as follows:
  - ▶ `EACCES`: no permission to operate
    - ★ tried `msgsnd()`, but no write permission
    - ★ tried `msgrcv()`, but no read permission
  - ▶ `EIDRM`: the message queue was removed (see later how to remove)
  - ▶ `EINTR`: the process caught a signal while waiting on a blocking `msgsnd()/msgrcv()` call
    - ★ on full queue for `msgsnd()`, or
    - ★ no selected message available for `msgrcv()`
    - ★ `msgsnd()` and `msgrcv()` **cannot** be restarted after the handler with flag `SA_RESTART`
  - ▶ `ENOMEM`, `E2BIG`: system limits reached

# Controlling (and deleting) a message queue by msgctl()

- The system call msgctl() enables several actions to be performed on the message queue

```
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
```

- ▶ `msqid`, is the ID of the queue
- ▶ `cmd`, describes the action to be taken over the queue
- ▶ `buf`, is a parameter for the action (see next for details)

- To remove and deallocate the queue `msqid`

```
int msgctl(int msqid, IPC_RMID, NULL);
```

- ▶ after the queue is removed, processes blocked on `msgrcv()` on the queue `msqid` will be unblocked with `errno=EIDRM`



# Controlling (and deleting) a message queue by msgctl()

- To get the status of the queue

```
int msgctl(int msqid, IPC_STAT, struct msqid_ds *buf);
```

it will return the data structure of the queue (man msgctl)

```
struct msqid_ds {  
    struct ipc_perm msg_perm; /* Owner, permission */  
    time_t msg_stime;         /* Time of last msgsnd */  
    time_t msg_rtime;         /* Time of last msgrcv */  
    time_t msg_ctime;         /* Time of last change */  
    msgqnum_t msg_qnum;       /* Cur # msg in queue */  
    msglen_t msg_qbytes; /* Max bytes allowed in Q */  
    pid_t msg_lspid;          /* PID of last msgsnd */  
    pid_t msg_lrpid;          /* PID of last msgrcv */  
};
```

## Queues vs. pipes

- Message queues offer an IPC facility similar to pipes

	<b>pipes</b>	<b>message queues</b>
unit of data	byte	message (any user-defined data structure)
terminology	write, read, file descriptors	send, receive, IDs
lifecycle	closed after all read/write file descriptors are closed	<i>persistent</i> : stay alive even after all processes (creator, senders, receivers) terminates
read blocks	if empty & some write ends are open	always if empty
write blocks	if full	if full
deallocation	implicitly after all fd are closed	must be made <b>explicitly</b> by the user
abstraction	low	high

## Example of usage of message queues

- Sender process sends a message of type `argv[1]` to a queue. The text of the message is read from `stdin`  
*test-ipc-msg-snd.c*
- Receiver process receives a message of type `argv[1]` and prints its content to `stdout`. If a “special” message of type `MSGTYPE_RM` is received (which is a user-defined macro in *ipc-msg-common.h*), then the message queue is erased  
*test-ipc-msg-rcv.c*
- they share a common header file  
*ipc-msg-common.h*

## Example 2 of usage of a message queue

- The parent process:
  - 1 Create a queue
  - 2 Forks NUM\_PROC sender child processes
  - 3 Forks a receiver process
  - 4 Waits for the sender processes to terminate
  - 5 Waits for the queue to be empty
  - 6 Deallocate the queue, waits for the receiver, then exit
- Each sender child process:
  - 1 Sends NUM\_MSG to the queue of type from 1 to NUM\_MSG
- The receiver process:
  - 1 Receives all messages from the queue and prints them

*test-ipc-msg-fork.c*

# Message queues: POSIX APIs

- For historical reasons, the course follows the System V API
- However, today the POSIX standard is dominant
- Here is a one slide overview
  - ▶ `man mq_overview` for an overview of POSIX message queues
  - ▶ Messages in POSIX queues have a *priority* (similar to SysV *type*)
  - ▶ `mq_open(...)`, `mq_close(...)` to create and close the queue
  - ▶ `mq_send(...)`, `mq_receive(...)` to send/receive messages
  - ▶ `mq_notify(...)` to enable asynchronous notification of message queue via `struct sigevent`
  - ▶ when interrupted by a handled signal, blocking system calls may be restarted if `SA_RESTART` flag is set