

Operating Systems Lab (C+Unix)

Enrico Bini

University of Turin

Outline

1 SysV: semaphores

Why semaphores?

- In concurrent programming (many processes running simultaneously), the output depends of the input **and on the scheduling decisions**
- synchronization primitives are used to constrain the possible schedules
- *semaphores* are **synchronization primitives**

How does a semaphore work?

- For any semaphore s , the kernel records a *value* denoted by $v(s)$ **always ≥ 0**
- The value $v(s)$ of a semaphore represents the number **of available accesses to the resource** protected by s

- Any process can perform the following actions on a semaphore s :

- ① **1 Initialize the value $v(s)$ with some integer a (number of allowed concurrent accesses to the resource)**
 - ★ $v(s) \leftarrow a$
- ② **2 Use, if available, the shared resource** protected by the semaphore s
 - ★ if $v(s)$ equals 0, block the process until $v(s) > 0$
 - ★ decrement $v(s)$ and use the resource
- ③ **3 Release a resource being used**
 - ★ increment $v(s)$ (it is never blocking)
- ④ **4 Wait until $v(s)$ equals zero.** A process waiting until $v(s)$ is zero can be used to
 - ★ “refill” the resource
 - ★ have many processes waiting for the same “green light”

Creating/accessing an System V semaphore

- System V implements an **array of semaphores**: each operation onto an array of semaphores is **atomic**
 - ▶ an array is useful for code fragments that need more than one resource
- The system call

```
int semget(key_t key, int nsems, int semflag);
```

returns the identifier of an array of `nsems` semaphores associated to `key`

- ▶ `semflag` is a list of ORed ("|") options including:
 - ★ read/write permissions (least significant 9 bits)
 - ★ `IPC_CREAT`:
 - (1) create a new semaphore associated to the `key`, if it doesn't exist
 - (2) return the existing semaphore associated to the `key`, if it exists
 - ★ `IPC_EXCL` (used only with `IPC_CREAT`): the call fails (with `errno=EEXIST`) if the semaphore exists

- When created semaphores **are not initialized** to any value: they must be explicitly initialized by `semctl(...)` (see later)
- Semaphores are persistent object: they will survive to the process death, they must be erased explicitly

Operation on a single semaphore

- Access/release of a semaphore are called **semaphore operations**
- An operation on **a single semaphore** (in an array) is described by a dedicated data structure sembuf

```
struct sembuf {  
    unsigned short sem_num;    /* sem number */  
    short          sem_op;     /* sem operation */  
    short          sem_flg;    /* flags */  
}
```

- An **array of operations** over the semaphore s_id are performed by

```
int semop(int s_id, struct sembuf * ops, size_t nops);
```

- ▶ ops, array of the operations
- ▶ **nops, number of the operations in ops** (\leq size of semaphores' array)
- ▶ **Notice:** the nops operations are made all together atomically
- ▶ The call blocks if **any** of the operations cannot be made

Semaphore operations to access/release a resource

```
struct sembuf {  
    unsigned short sem_num;    /* sem number */  
    short          sem_op;     /* sem operation */  
    short          sem_flg;    /* flags */  
}
```

- To access a resource protected by semaphore the value of the semaphore must be **decremented** by setting `my_op.sem_op = -<num-res>;`
 - ▶ Important: the process **blocks** if <num-res> resources are not available
- To release a resource protected by semaphore the value of the semaphore must be **incremented** by setting `my_op.sem_op = <num-res>;`
 - ▶ Important: the process **never blocks** when increasing the resources
- `sem_num`, indicates the index of the semaphore in the array

Basic usage: protecting a critical section

- Example: more processes need to write to the same shared memory area
- To avoid that inconsistent data, **only one process is allowed in the critical section** of code modifying the data

```
struct sembuf my_op;
int sem_id;
...
sem_id = semget(IPC_PRIVATE /*key*/, 1 /*nsems*/, 0600 /*flags*/);
/* Initialize the semaphore allowing one access only */
semctl(sem_id, 0, SETVAL , 1); /* later details of semctl call */
/* sharing sem_id with all processes accessing the resource */
...
/* now trying to access critical section */
my_op.sem_num = 0; /* only one semaphore in array of semaphores */
my_op.sem_flg = 0; /* no flag: default behavior */
my_op.sem_op = -1; /* accessing the resource */
semop(sem_id , &my_op, 1); /* blocking if others hold resource */
/* NOW IN CRITICAL SECTION */
...
my_op.sem_op = 1; /* releasing the resource */
semop(sem_id , &my_op, 1); /* may un-block others */
```


Sem. op. to wait until semaphore is zero

```
struct sembuf {  
    unsigned short sem_num;    /* sem number */  
    short          sem_op;     /* sem operation */  
    short          sem_flg;     /* flags */  
}
```

- If processes A_1, A_2, \dots, A_n must wait that another process B reaches some given point, then
 - ① a semaphore is initialized with value 1
 - ② all processes A_1, A_2, \dots, A_n “waits for zero” with a semaphore operation
`my_op.sem_op = 0;`
 - ③ process B decrements the same semaphore by one by
`my_op.sem_op = -1;`
 - ★ the value of the semaphore becomes zero and the processes A_1, A_2, \dots, A_n will be unblocked

Don't wait forever

- If a resource protected by a semaphore is unavailable, a process may:
 - ① wait until the resource is available again (as seen before), or
 - ② decide to do something else
- The flag `IPC_NOWAIT` may be set in a semaphore operation

```
struct sembuf sop;  
sop.sem_flg = IPC_NOWAIT;  
semop(.., &sop, 1); /*dont wait*/
```

- When executing the `semop()`
 - ▶ if the resource is available, get it as usual
 - ▶ if unavailable don't wait, return `-1`, and `errno` set to `EAGAIN`
- If only waiting for some time is desired

```
#include <time.h>  
struct timespec {  
    time_t tv_sec; /* seconds */  
    long tv_nsec; /* nanoseconds */ }  
semtimedop(/*same as semop*/, struct timespec * timeout);
```

timeout



Controlling (and initializing) a semaphore

- The system call `semctl()` enables several actions to be performed on semaphores

```
int semctl(int s_id, int i, int cmd);  
int semctl(int s_id, int i, int cmd, /* arg */);
```

- ▶ `s_id`, is the ID of the semaphore set
 - ▶ `i`, is the index of the semaphore in the set
 - ▶ `cmd`, describes the action to be taken over the semaphore
 - ▶ the optional fourth argument depends on the type of command
- To set (initialize) or get the value of the `i`-th semaphore in a set

```
int semctl(int s_id, int i, SETVAL, int val);  
int semctl(int s_id, int i, GETVAL);
```

- ▶ if `GETVAL`, the value of the `i`-th semaphore is returned;

Semaphore: getting information, removing

- To know how many processes are **blocked**

```
int semctl(int s_id, int i, GETNCNT);
```

returns the number of processes waiting for the i-th semaphore to increase

- To know the process who **last accessed a** resource

```
int semctl(int s_id, int i, GETPID);
```

returns the PID of the last processes who executed a semop(s_id, ...) operation on the i-th semaphore

- To deallocate the semaphore s_id

```
int semctl(int s_id, /*ignored*/, IPC_RMID);
```

- ▶ when a process is blocked on a semop(id, ...) and the semaphore is removed by semctl(id, ..., IPC_RMID) the process is unblocked with return value -1 and errno is set to EIDRM

Semaphores and signals


- When a process is blocked on a `semop(...)` and it receives a non-masked signal
 - ① the handler is executed
 - ② the `semop()` system call returns `-1` and `errno` is set to `EINTR`
- Even if the flag `SA_RESTART` was set in the signal handler by the `sigaction(...)`, an interrupted `semop(...)` will always fail with `errno` set to `EINTR`

Wrong ways to wait for a semaphore

- Do not loop forever testing the value of a semaphore

```
sop.sem_flg = IPC_NOWAIT;  
do {  
    semop(.., &sop, 1);  
    while (errno == EAGAIN);
```

Semaphores: Examples



- 1 Processi figli che scrivono nella pipe in modo ordinato
test-pipe-round.c
which uses a small module for handling semaphores
 - ▶ *my_sem_lib.h* (header file)
 - ▶ *my_sem_lib.c* (implementation of the functions)
- 2 Tanti processi che vogliono cucinare condividendo le risorse di una cucina
test-sem-cook.c

Semaphores: POSIX APIs

- For historical reasons, the course follows the System V API
- However, today the POSIX standard is dominant
- Here is a one slide overview
 - ▶ `man sem_overview` for an overview of POSIX semaphores
 - ▶ `sem_open(...)`, `sem_init(...)`, and `sem_destroy(...)` to create, initialize and destroy a semaphore
 - ▶ `sem_post(...)` to increment by one, `sem_wait(...)` to decrement by one. `sem_timedwait(...)` to wait at most a given timeout
 - ▶ no “wait-for-zero” interface
 - ▶ not possible to increment by more than one
 - ▶ no array of semaphore operations
 - ▶ simpler interface (which is good and bad)