



# Operating Systems Lab (C+Unix)

**Enrico Bini**

University of Turin

# Outline

- 1 SysV: shared memory

## IPC shared memory

- System V implements *shared memory*: remember, when allocating by `malloc` you are allocating over the heap (which is private to the process!)
  - ▶ If memory is allocated by `malloc` and then a process is forked
  - ▶ the two processes **do not share** the allocated memory
  - ▶ pointers represent logical addresses (not physical)
- **Shared memory is a fast way for processes** to communicate: no kernel structure (buffers, queues, etc) mediating the access to the shared memory:
  - ▶ **this is good: fast way to implement IPC**
  - ▶ **this is bad: high risk of inconsistent** behavior if memory is read “in the middle” of a write
- Once a process writes to a shared memory segment, the data written becomes **immediately accessible** to the other processes accessing the shared memory segment
- **Simplicity of usage: assignments are made with the same syntax of private memory**: no special function to access, no write, read, `msgsnd`, `msgrcv`, just “=”

# Lifecycle of a shared memory segment

- 1 **Creation** by `shmget(...)`: **size** of memory must be specified
  - ▶ the shared memory segment is allocated over an area accessible to all processes
- 2 **Attaching** the shared memory area to the **process address space**
  - ▶ **after** a shared segment is attached to a private address space, it **can be normally used by the process**
  - ▶ any data written to the shared memory segment becomes immediately visible to other processes sharing the same segment
- 3 **Detaching** the segment from the process address space
  - ▶ the shared memory is **no longer visible, but it still exists (it is a persistent object)**
- 4 **Deallocation** of the shared memory segment

# Creating a shared memory segment

- The system call

```
int shmget(key_t key, size_t size, int shmflg)
```

returns the identifier of a shared memory segment associated to key of size **at least** size (the allocated size is a multiple of PAGE\_SIZE)

- ▶ shmflag is a list of ORed ("|") options including:
  - ★ read/write permissions (least significant 9 bits)
  - ★ IPC\_CREAT:
    - (1) create a new shm segment associated to the key, if it doesn't exist
    - (2) return the existing shm ID associated to the key, if it exists
  - ★ IPC\_EXCL (used only with IPC\_CREAT): the call fails (with errno=EEXIST) if the shm segment exists
- Shared memory segments are persistent object: they will survive to the process death, they must be erased explicitly

# Attaching a shared memory area to a process

- To attach a shared memory segment to the address space of a process

```
void *shmat(int shmid, NULL, int shmflg)
```

- ▶ `shmid`, ID of the shm object
- ▶ second argument used for advanced features: setting to NULL is safe
- ▶ `shmflg`, flags
  - ★ `SHM_RDONLY`, uses the shared memory in read-only mode
  - ★ plus others for advanced settings
- it returns a pointer to the shared memory segment
- Typical usage (malloc-like)

```
struct my_data * datap; /* shared data struct */  
  
shm_id = shmget(IPC_PRIVATE, sizeof(struct my_data), 0600);  
datap = shmat(shm_id, NULL, 0);  
/* From now on, all processes accessing to  
   datap->something, read/write in shared mem */
```

# Detaching a shared memory

- A shm segment is detached by

```
int shmdt(const void *shmaddr);
```

- shmaddr is the address of the segment we want to detach, previously returned by a shmat call
- **Implicit detaching** of a shm segment occurs when:
  - 1 the process terminates
  - 2 the control flow passes to another process by exec()
- detaching is **not** deallocation

# Control (and deallocation) of a shm segment

- A shared memory segment is controlled by

```
int shmctl(int shmid, int cmd,  
           struct shmid_ds * buf);
```

- ▶ shmid, the ID of the shared memory object
- ▶ cmd, is the command to be made (IPC\_STAT, IPC\_RMID, ...)
- ▶ the third argument may be used depending on the command cmd

- To mark a shared memory for deallocation

```
int shmctl(int shmid, IPC_RMID, NULL);
```

- ▶ **Important:** the actual deallocation happens only when the last process is detached from the shared memory segment
- ▶ Deallocating the shm segment immediately would create problems to the processes still using the segment
  - ★ these problems cannot be detected by some `errno` (as for message queues), because the access to memory segment is made by assignments "=", not by any function calls



## Example on shared memory

- Many child processes filling a shared table
- Each process needs to get a unique entry in the table, then it can write without conflict
  - ▶ *Makefile*
  - ▶ *test-shm.h*
  - ▶ *test-shm-parent.c*
  - ▶ *test-shm-child.c*

# Shared memory: POSIX APIs

- For historical reasons, the course follows the System V API
- However, today the POSIX standard is dominant
- Here is a one slide overview
  - ▶ `man shm_overview` for an overview of POSIX shared memory
  - ▶ POSIX API uses file descriptor instead of IDs
  - ▶ `shm_open(...)` similar to `shmget(...)`
  - ▶ `mmap(...)` similar to `shmat(...)`, but can do much more (mapping any file descriptor to memory space)
  - ▶ `shm_unlink(...)` similar to `shmdt(...)`